

Intro to **Machine Learning**

Vinayaka Gude, Ph.D.
Elon University

Agenda

01

Artificial Intelligence

02

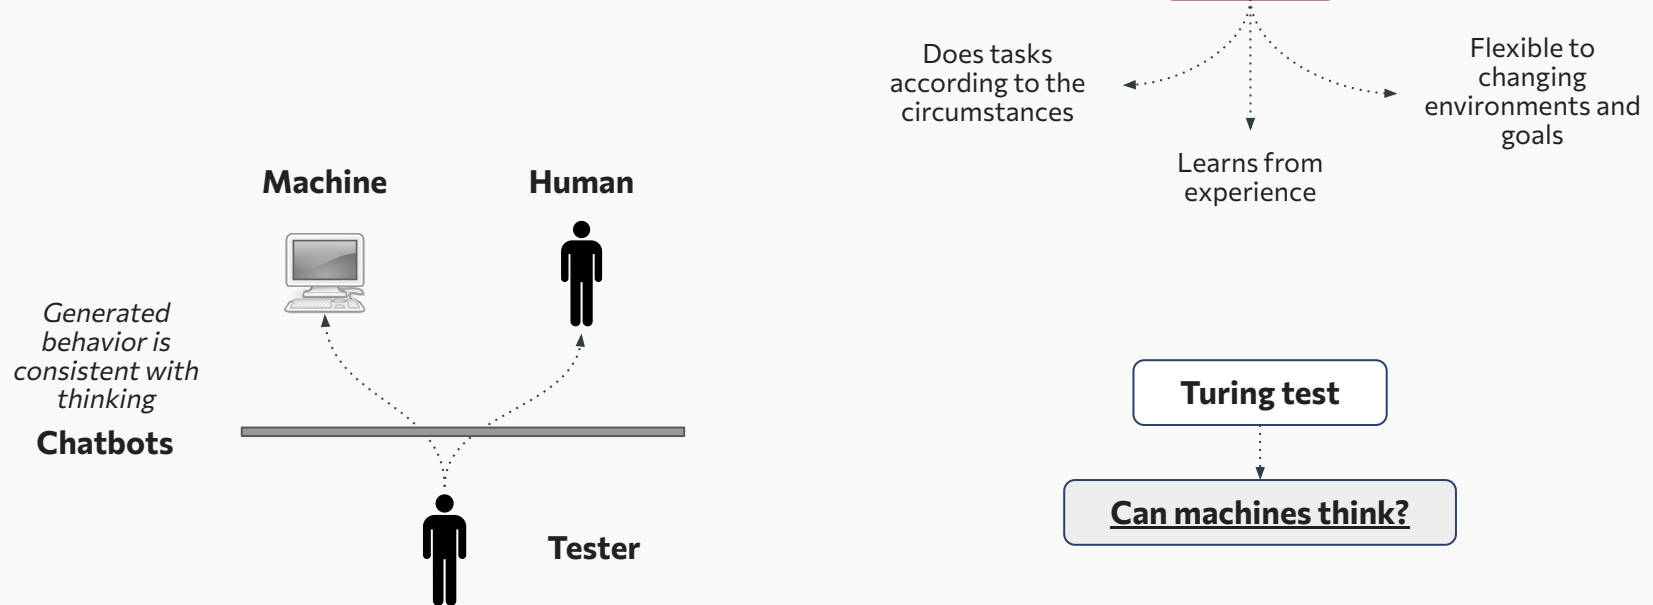
Introduction to Machine Learning

03



ML Challenges

Artificial Intelligence

The study of computer systems that attempt to model and apply the **intelligence** of the human mind.

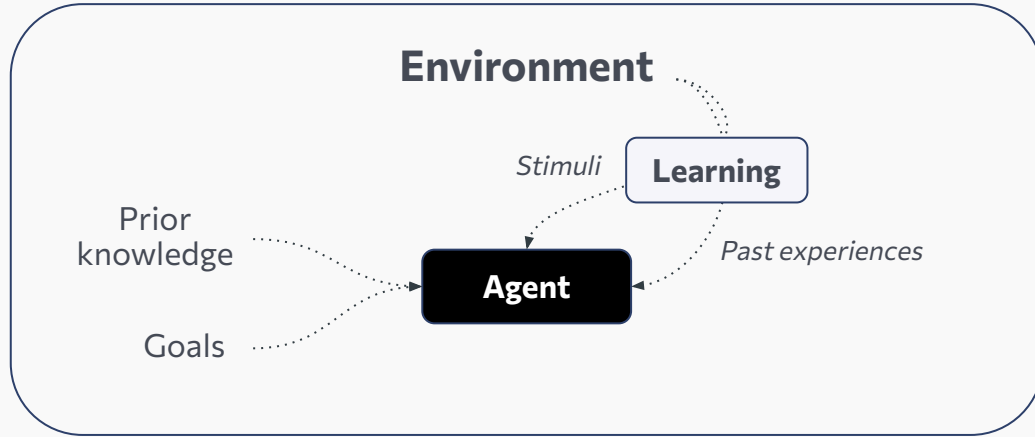


Human Intelligence & Beyond



Ant colonies

Artificial Intelligence



Intelligence → **General** intelligence

Specialized intelligence

Humans + computers → **Collective** intelligence


Why do we need intelligence?

Data → Knowledge

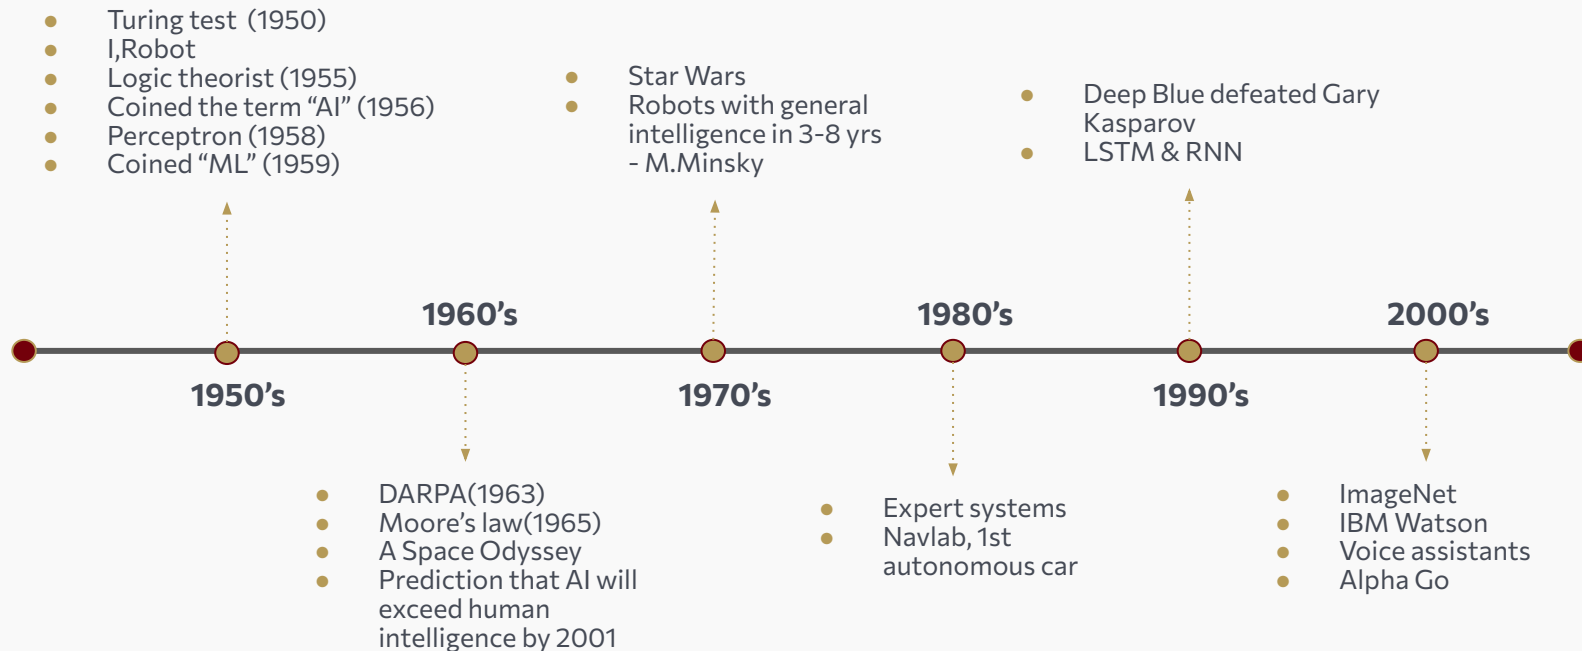
Avoid providing **manual rules** and going through **huge amount of data**

Recent developments:

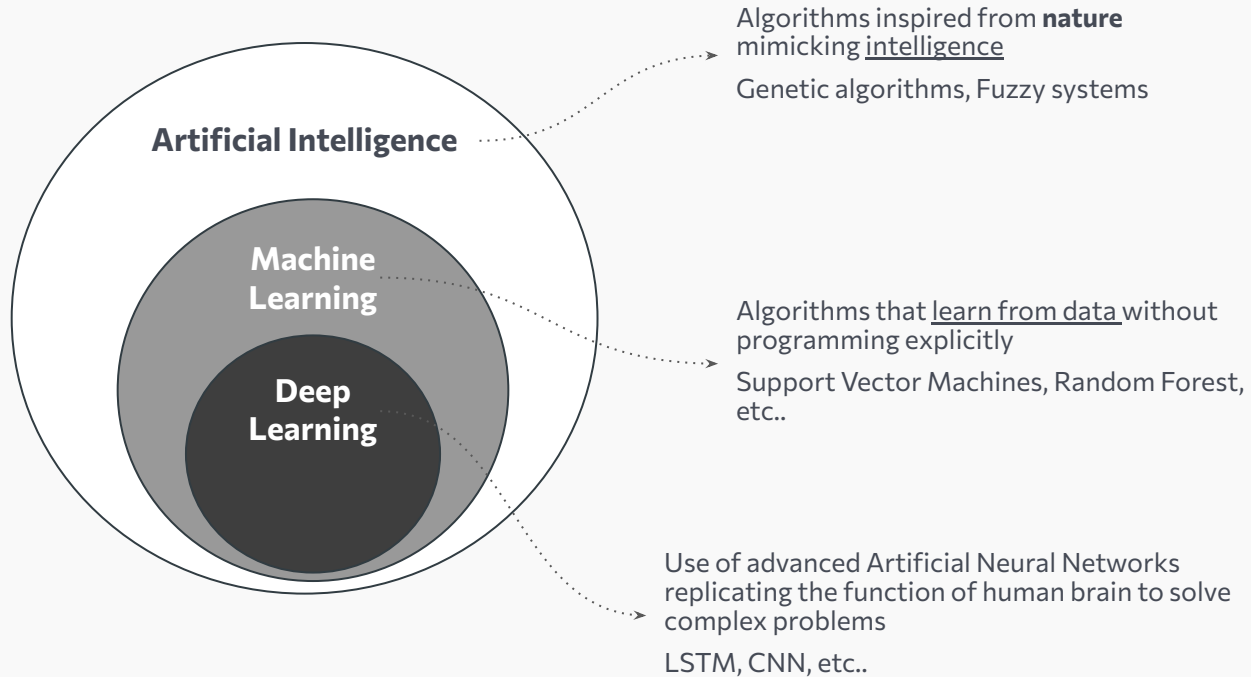
- Detecting Skin cancer
- Investment Management
- AI powered logistics
- Speller100



Brief History of AI



AI/ML/DL



AI applications

INPUT	OUTPUT	APPLICATION
Voice Recording	Transcript	Speech Recognition
Image	Caption	Image Recognition
Videos/movies/posts	Relevant content	Recommendation Systems
Customer Data	Ads	Targeting Ads
Historical Market Data	Future Market Data	Trading Bots
Transaction details	Fraud labels	Fraud detection
Purchase history	Future purchase behavior	Customer retention
Faces	Names	Face Recognition
Car locations and speed	Traffic Flow	Traffic Lights Automation

Data is the new oil

Data → Better Decisions → Profits


Machine Learning

Science of programming computers so they can learn from data.

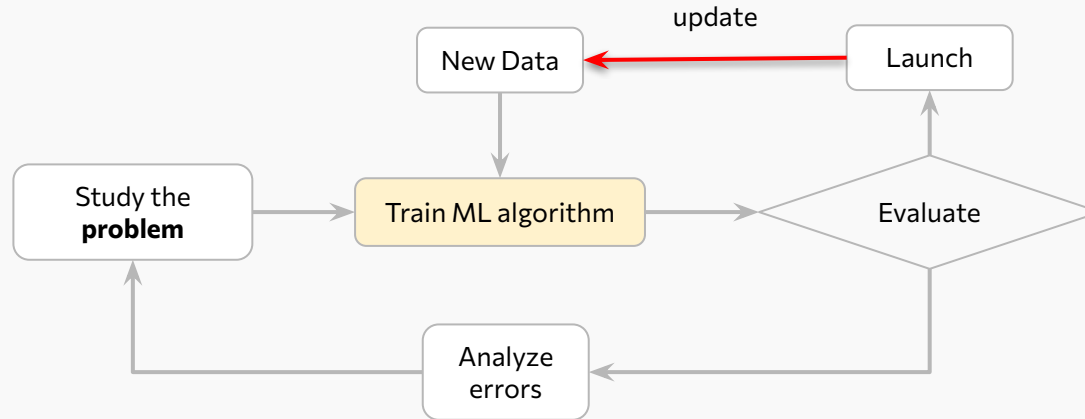
Machine Learning is the field of study that gives computers the **ability to learn** without being explicitly programmed - *Arthur Samuel, 1959*

A computer program is said to learn from **experience** (E) with respect to some **task** (T) and some **performance measure** (P), if its performance on (T), as measured by (P), improves with experience (E). - *Tom Mitchell, 1997*


Why use Machine Learning?



ML approach

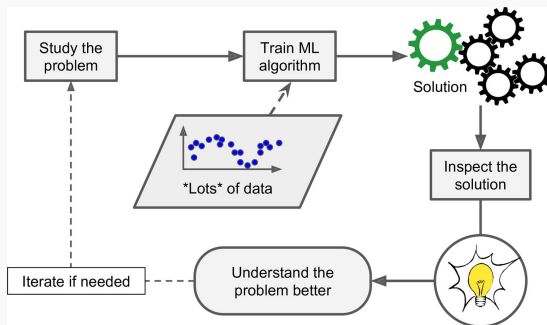


How does ML work?



Output = **Function** (data)

Capable of dealing with complex problems such as **speech recognition, image/video analysis** and **vast amount of numerical data**.




Helps us discover new patterns → **Data Mining**

Can **adapt** to new data

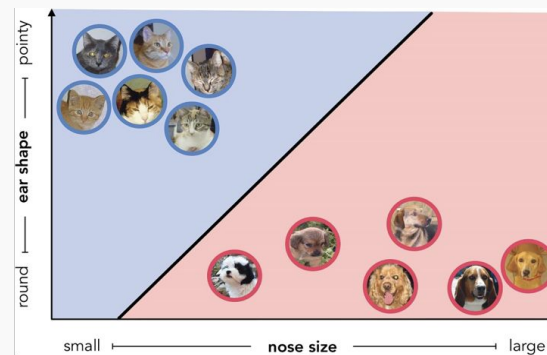
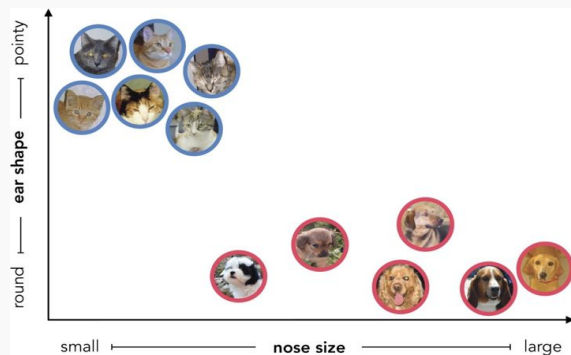


Types of **Learning**

Supervised Learning





Examples:
**Classification &
Regression**



Training data includes labels


Supervised Learning



Linear Regression, Logistic Regression, Random Forest, Support Vector Machine & Neural Networks

Logistic Regression is used for **binary classification**

Unsupervised Learning



Dimensionality Reduction

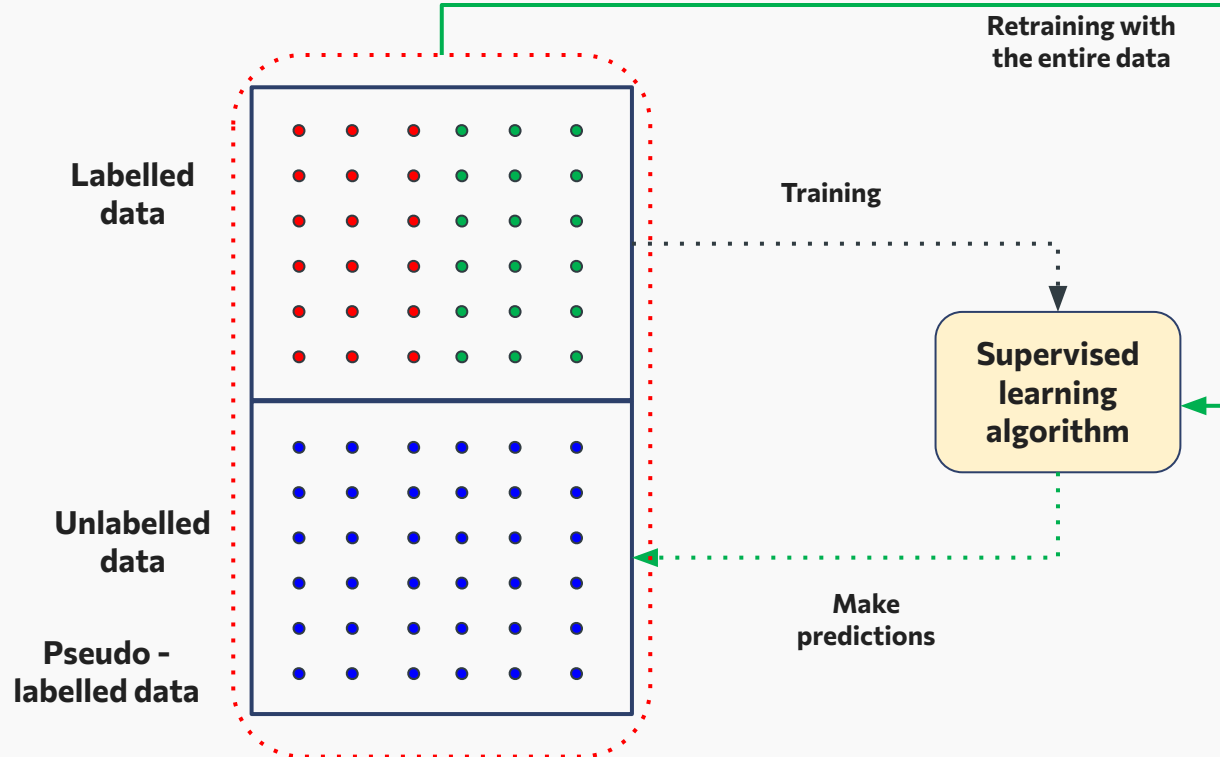
Simplify data without losing too much information

Principal Component Analysis (PCA)


Association rules

Apriori

Semi-Supervised Learning



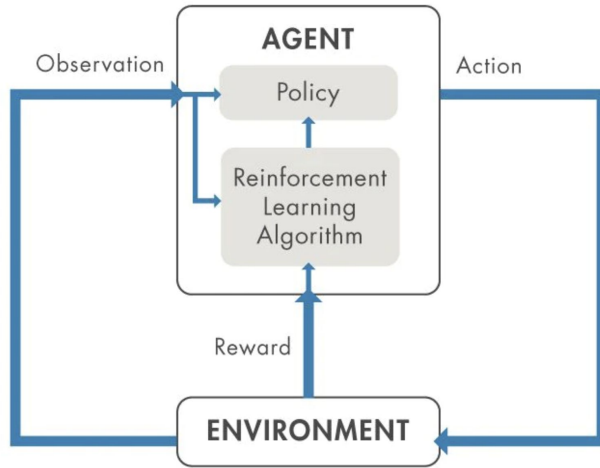
Semi-Supervised Learning



Semi-supervised learning models are a **combination** of supervised and unsupervised models


Labeling data can be **costly & time consuming**

Reinforcement Learning




Algorithm → updates the policy
based on the reward

Autonomous driving



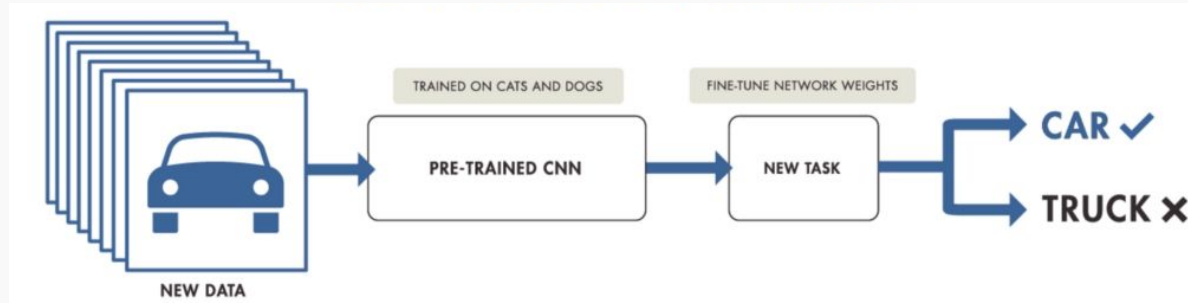
- 1 Observe
- 2 Select action using policy



- 3 Action!
- 4 Get reward or penalty



Transfer Learning



Saves times - no need to build models from scratch.

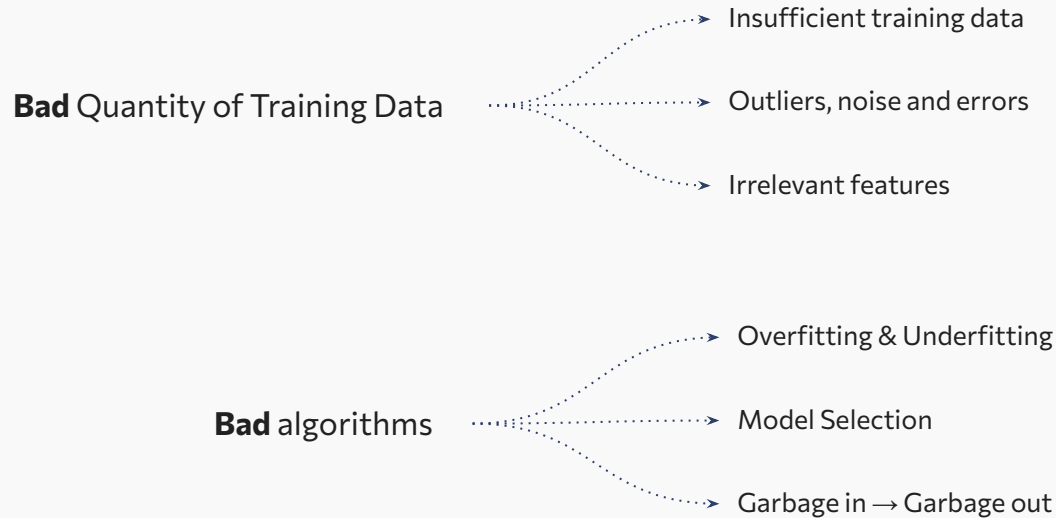
Eliminates the need for **huge training dataset**

Improves **generalizability**

Which ML approach would you use?


- Develop a model that integrates with wearable devices to help patients manage chronic diseases like diabetes or hypertension. The agent should provide personalized recommendations for medication, diet, and exercise based on continuous monitoring of the patient's physiological data.
- A hospital wants to predict how long a patient is likely to stay based on their diagnosis, medical history, and current condition.
- An insurance company wants to identify potentially fraudulent claims to reduce losses and prevent abuse.
- A public health organization wants to design targeted campaigns to promote healthy lifestyles. They need to segment the population into groups with shared characteristics.

ML - Challenges




No Free Lunch



Effectiveness of Data



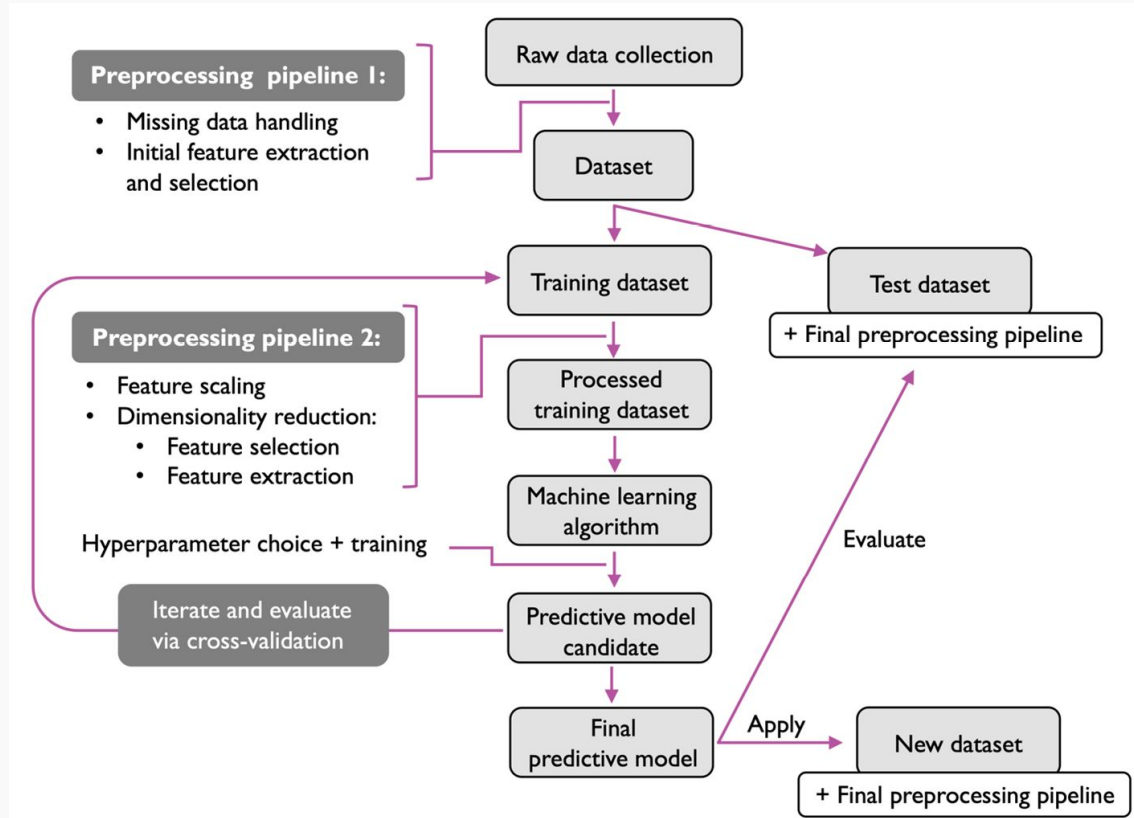
Non-representative training data





Overfitting



ML Process



Overfitting



Keeping up with AI research

INPUT	OUTPUT
New code (Research)	https://paperswithcode.com/
New models (Open source)	https://huggingface.co/models?sort=trending
News (HBR, MIT)	https://www.technologyreview.com/
Conference Proceedings	https://aaai.org/aaai-publications/aaai-conference-proceedings/
Podcast	https://emerj.com/artificial-intelligence-podcast/
LinkedIn	https://www.linkedin.com/in/yann-lecun/

Podcast Task

Pick an episode of your choice from the **year 2024** from "**AI in Business**" podcast. Submit a summary of all three episodes after going through them. The summary should contain the problem, challenges, and AI solution.

<https://podcasts.apple.com/us/podcast/the-ai-in-business-podcast/id670771965>

Thank you!

Any questions?

gude.vinayaka@outlook.com